


Tree Traversals

- Print the content of a BST in sorted order => inorder traversal

```
void printTree( ostream & out = cout ) const
{
    if( isEmpty( ) )
        cout << "Empty tree" << endl;
    else
        printTree( root, out );
}
```

```
/* Internal method */
void printTree( BinaryNode *t, ostream & out ) const
{
    if( t != NULL )
    {
        printTree( t->left, out );
        cout << t->element << endl;
        printTree( t->right, out );
    }
}
```

$T(N) = ?$




Tree Traversals

- Compute the height of a tree
=> postorder traversal

```
/* Internal method */
int height( BinaryNode *t)
{
    if( t == NULL )
        return -1;
    else
        return 1 + max( height(t -> left), height( t ->right ))
}

T(N) = ?
```



B - Trees


- We assumed that the entire data structure fits in the main memory.
- What if some of the data resides on disk ?
=> Big-Oh does not work
- Big-Oh assumes that all operations take equal amount of time.
- Example:
~100,000 MIPS CPU (Intel Core i7, 4C, 3.4 GHz)
=> 100 billion instructions per second
Disk: 7200 rpm => one revolution per 1/120 sec. => one revolution takes 8.3 ms => average time = 8.3 ms
=> 120 disk accesses/sec.
One disk access => ~1 billion instructions
The number of accesses will dominate the running time.

Typical Search Tree on Disk

- Florida driving records = 10,000,000 items
- Each key (name) = 32 B, record = 256 B
- Approx 2.3 GB, assume that it does not fit in memory.
- There are 20 users in the system \Rightarrow $1/20$ of resources
 \Rightarrow in 1 sec: 5 billion instructions or 6 disk accesses.
- **Unbalanced BST:** 10,000,000 disk accesses \Rightarrow 1.6 million sec. (about 19 days!)
- **Average BST:** $1.38 \log N = 32$ disk accesses \Rightarrow 5 sec.
- **AVL Tree:** average case $\log N = 25$ disk accesses \Rightarrow 4 sec.
- **Problem:** We cannot go below $\log N$ using BST!
- **Solution ?** \Rightarrow more branching reduces height

M-ary Search Tree

- Allows M-way branching \Rightarrow height = $\log_M N$
- Example: 5-ary tree of 31 nodes
- BST of 31 nodes \Rightarrow 5 levels
- 5-ary tree \Rightarrow 3 levels
- BST: we need one key to decide which child to explore
- M-ary Tree: we need M-1 keys to decide which child to explore



B - Trees

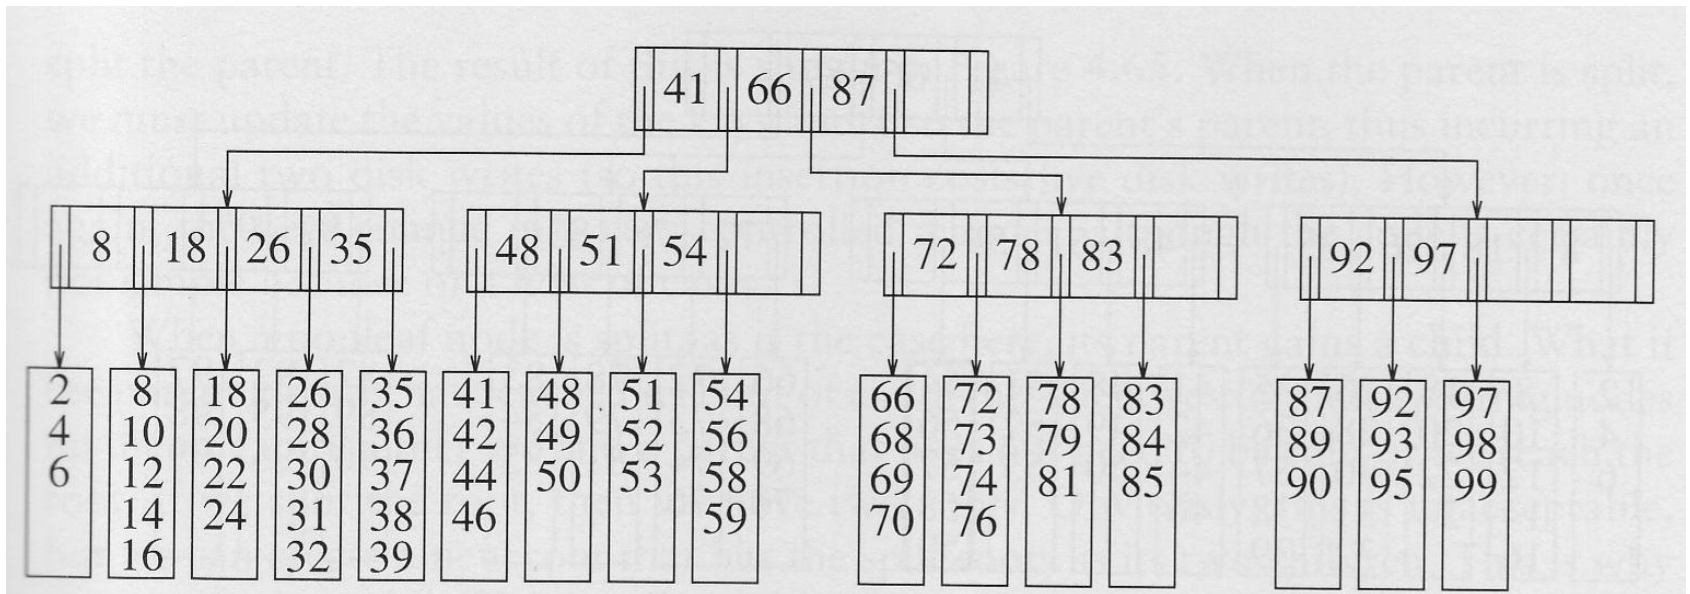
R. Bayer and E. McCreight (1972):

"Organization and Maintenance of Large Ordered Indexes", Acta Informatica 1 (3):
173–189, 1972

A B – Tree of order M is an M-ary tree with the following properties:

- The data items are stored at leaves.
- The non-leaf nodes store up to M-1 keys to guide the searching; key i represents the smallest key in subtree i+1.
- The root is either a leaf or has between 2 and M children.
- All non-leaf nodes (except the root) have between $\lceil M/2 \rceil$ and M children.
- All leaves are at the same depth and have between $\lceil L/2 \rceil$ and L data items, for some L.

B – Tree Example



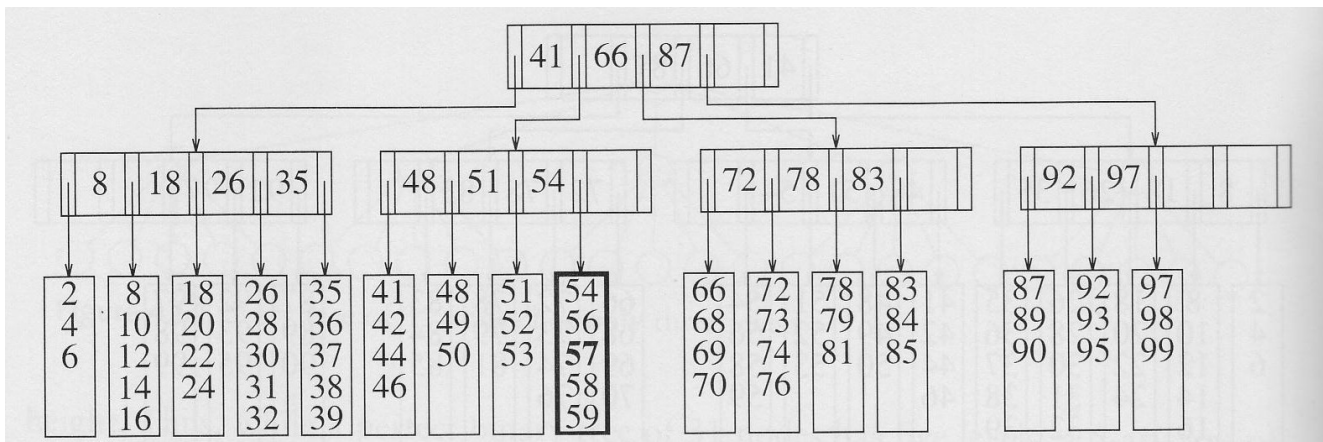
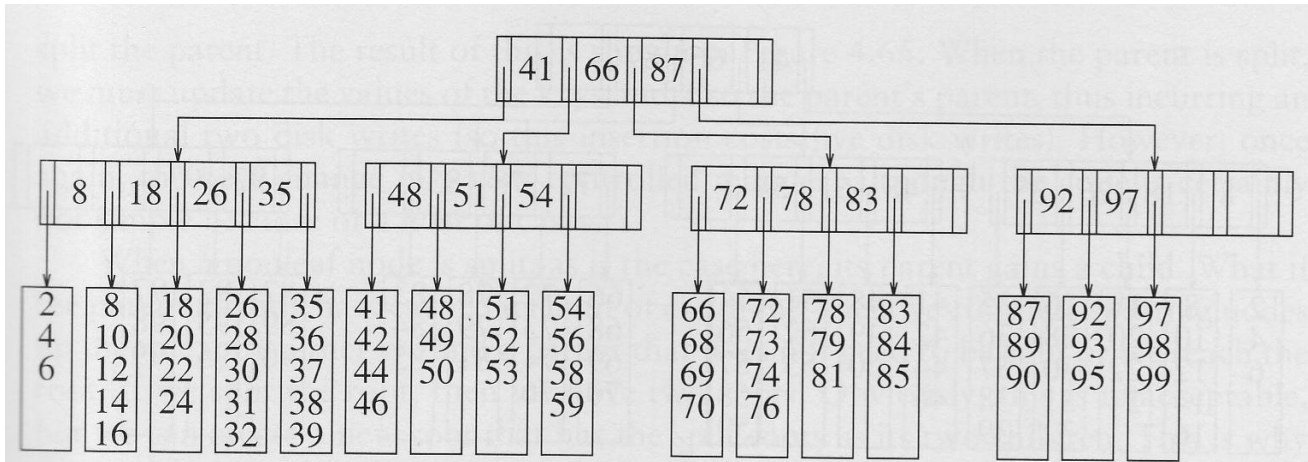
- $M = 5, L = 5$
- Nodes need to be half full => guarantees that the B-tree does not degenerate into a binary tree.

B - Trees

- Each node = disk block
- How to select M and L?
- Example:
 - Disk block = 8192 Bytes
 - key = 32 Bytes \Rightarrow Size of nonleaf node = $32(M-1) + 4M = 36M - 32$ Bytes
 - Pick largest M such that $36M - 32 \leq 8192 \Rightarrow M = 228$
 - Data record = 256 Bytes \Rightarrow 32 records/block $\Rightarrow L = 32$.
 - Each leaf has between 16 and 32 records \Rightarrow at most 625,000 leaves.
 - Leaves are at level 4 (worst case) \Rightarrow 3 disk accesses if we cache the root.
- Worst-case number of accesses = $O(\log_{M/2} N)$

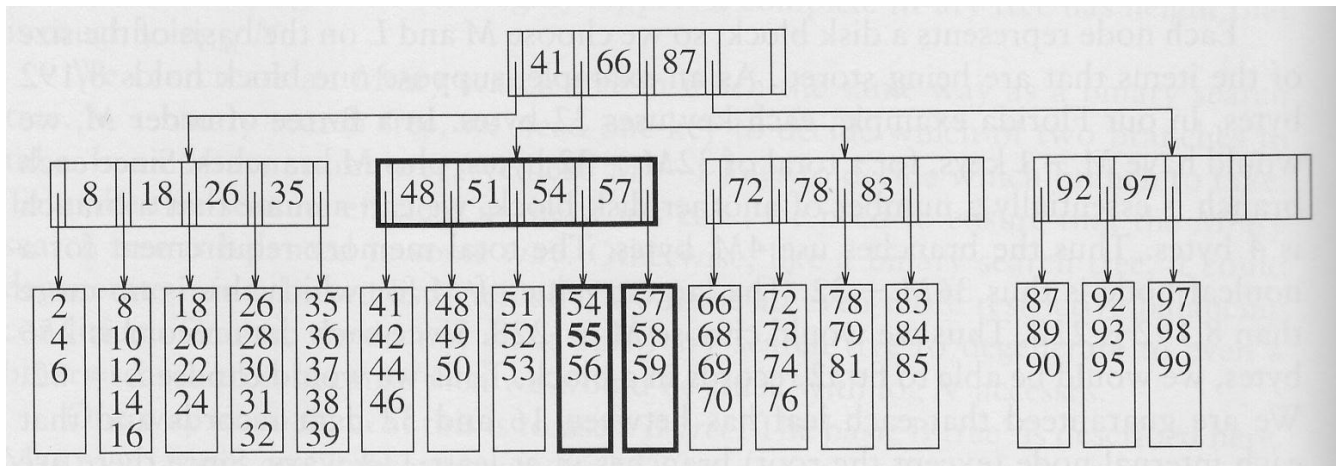
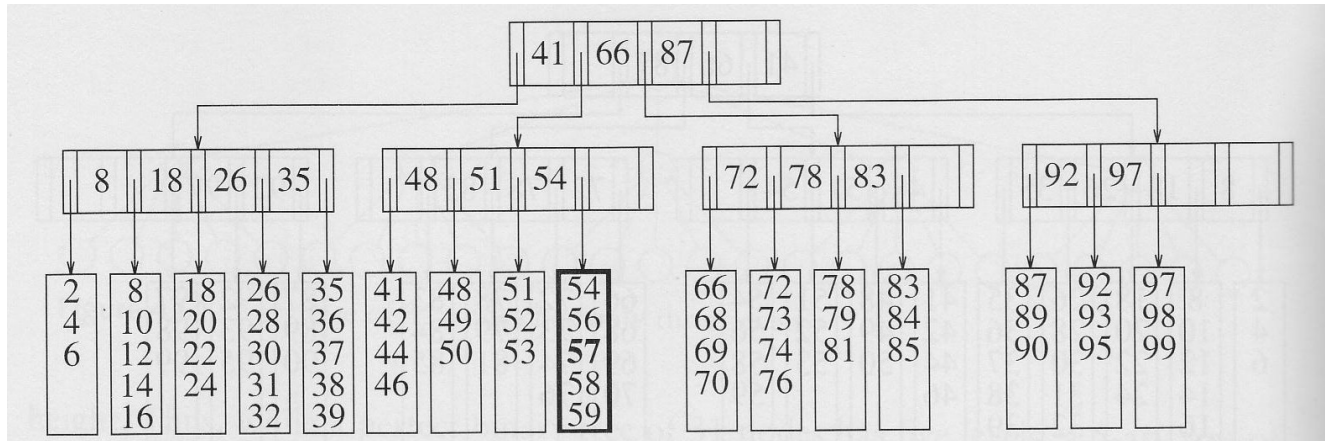
Insertion

Insert 57



Insertion

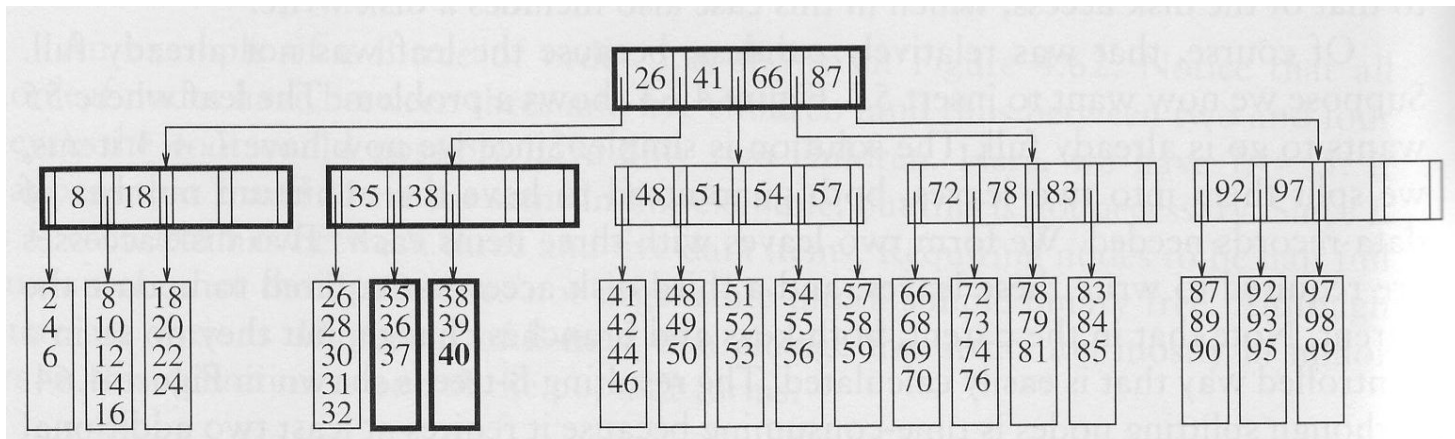
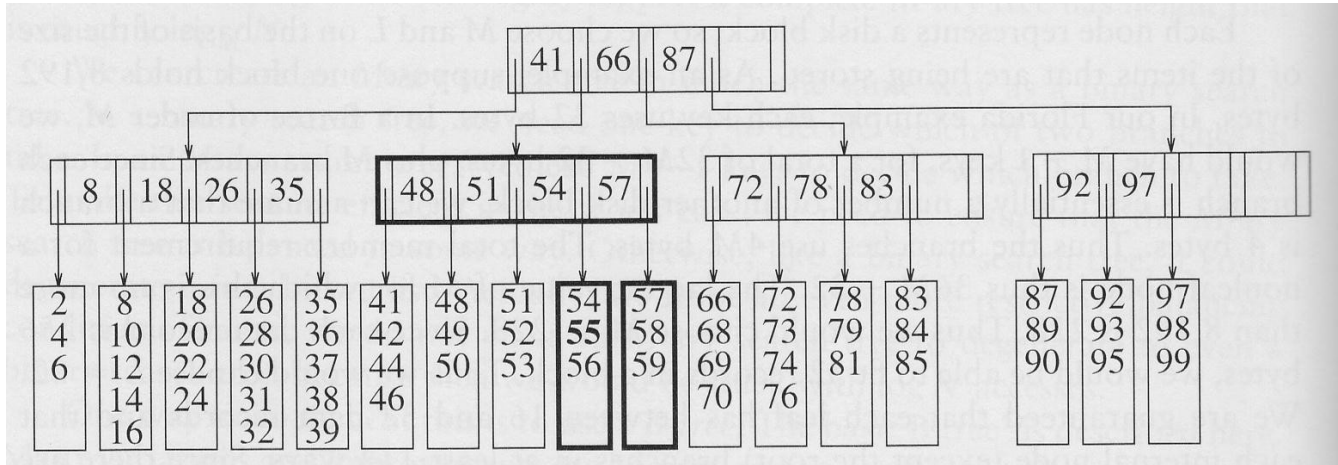
Insert 55



3 additional accesses for updating
Splitting is rare.

Insertion


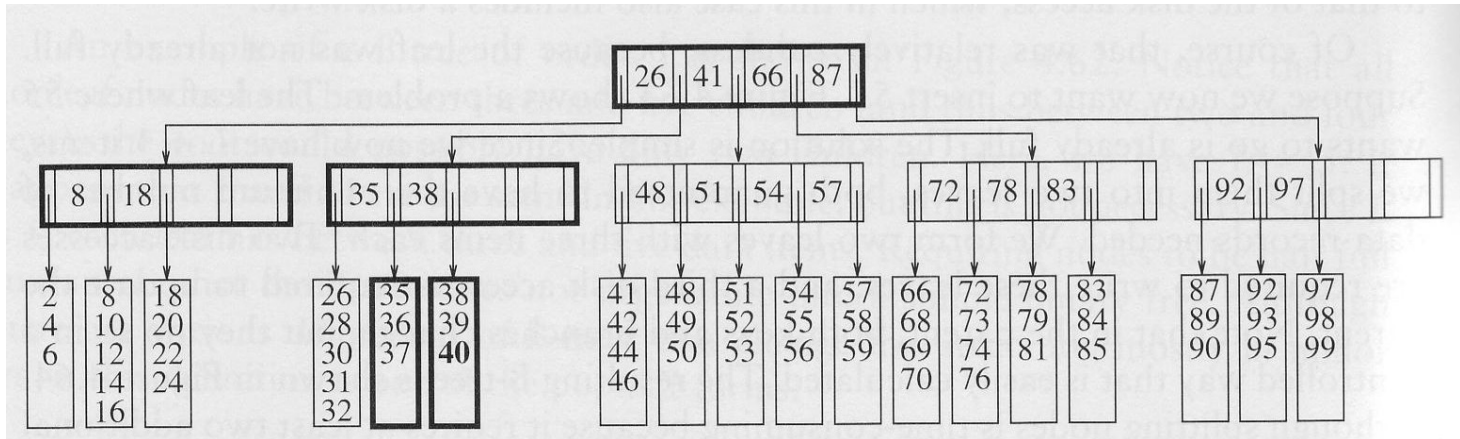
Insert 40



5 additional accesses for updating

Deletion

Delete 99



4 additional accesses for updating

Tree Applications

- Operating systems (file systems)
- Compilers (parse trees)
- Data bases (search trees)
- **Sorting algorithm:** insert items in a BST and then perform an inorder traversal $\Rightarrow O(N \log N)$